

nonlinear_regression

July 22, 2020

1 Nonlinear regression exercises


1.1 Linear regression review.

```
[1]: # Import the necessary libraries.  
import numpy as np  
import matplotlib.pyplot as plt
```

```
[2]: # Create data.  
data_size = 20  
X = np.linspace(0, 1, data_size)  
# Include noise.  
noise = np.random.uniform(low=-1.0, high=1.0, size=data_size) * 0.2  
y = np.sin(2.0 * np.pi * X) + noise
```

```
[3]: X_line = np.linspace(0, 1, 1000)  
sin_X = np.sin(2.0 * np.pi * X_line)
```

```
[4]: # Drawing functions for training data and correct answer data.  
def plot_sin():  
    plt.scatter(X, y)  
    plt.plot(X_line, sin_X, 'red')  
plot_sin()
```




```
[5]: from sklearn.linear_model import LinearRegression
```

```
[6]: # Predicted by linear regression.  
lin_reg = LinearRegression().fit(X.reshape(-1, 1), y)
```

```
[7]: # Check the regression variables and intercepts.  
lin_reg.intercept_, lin_reg.coef_
```

```
[7]: (0.8284329606722655, array([-1.63413944]))
```

```
[8]: # Draw the model.  
plt.plot(X_line, lin_reg.intercept_ + lin_reg.coef_ * X_line)  
plot_sin()
```



1.2 polynomial regression

```
[9]: # Predict the model by the square.
X_2 = X ** 2
```

```
[10]: X_new = np.concatenate([X.reshape(-1, 1), X_2.reshape(-1, 1)], axis=1)
```

```
[11]: X_2
```


```
[11]: array([0.          , 0.00277008, 0.01108033, 0.02493075, 0.04432133,
          0.06925208, 0.09972299, 0.13573407, 0.17728532, 0.22437673,
          0.27700831, 0.33518006, 0.39889197, 0.46814404, 0.54293629,
          0.6232687 , 0.70914127, 0.80055402, 0.89750693, 1.          ])
```

```
[12]: lin_reg_2 = LinearRegression().fit(X_new, y)
```

```
[13]: lin_reg_2.intercept_, lin_reg_2.coef_
```

```
[13]: (0.7786933756057179, array([-1.31912206, -0.31501737]))
```

```
[14]: plt.plot(X_line, lin_reg_2.intercept_ + lin_reg_2.coef_[0] * X_line + lin_reg_2.
      ↪coef_[1] * X_line ** 2)
plot_sin()
```



```
[15]: from sklearn.preprocessing import PolynomialFeatures
```

```
[16]: # Polynomial regression.
poly = PolynomialFeatures(degree=3)
poly.fit(X.reshape(-1, 1))
X_poly_3 = poly.transform(X.reshape(-1, 1))
```


```
[17]: X_poly_3
```

```
[17]: array([[1.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00],
 [1.00000000e+00, 5.26315789e-02, 2.77008310e-03, 1.45793847e-04],
 [1.00000000e+00, 1.05263158e-01, 1.10803324e-02, 1.16635078e-03],
 [1.00000000e+00, 1.57894737e-01, 2.49307479e-02, 3.93643388e-03],
 [1.00000000e+00, 2.10526316e-01, 4.43213296e-02, 9.33080624e-03],
 [1.00000000e+00, 2.63157895e-01, 6.92520776e-02, 1.82242309e-02],
 [1.00000000e+00, 3.15789474e-01, 9.97229917e-02, 3.14914711e-02],
 [1.00000000e+00, 3.68421053e-01, 1.35734072e-01, 5.00072897e-02],
 [1.00000000e+00, 4.21052632e-01, 1.77285319e-01, 7.46464499e-02],
 [1.00000000e+00, 4.73684211e-01, 2.24376731e-01, 1.06283715e-01],
 [1.00000000e+00, 5.26315789e-01, 2.77008310e-01, 1.45793847e-01],
 [1.00000000e+00, 5.78947368e-01, 3.35180055e-01, 1.94051611e-01],
 [1.00000000e+00, 6.31578947e-01, 3.98891967e-01, 2.51931768e-01],
 [1.00000000e+00, 6.84210526e-01, 4.68144044e-01, 3.20309083e-01],
 [1.00000000e+00, 7.36842105e-01, 5.42936288e-01, 4.00058318e-01],
 [1.00000000e+00, 7.89473684e-01, 6.23268698e-01, 4.92054235e-01],
 [1.00000000e+00, 8.42105263e-01, 7.09141274e-01, 5.97171599e-01],
```

```
[1.00000000e+00, 8.94736842e-01, 8.00554017e-01, 7.16285173e-01],
[1.00000000e+00, 9.47368421e-01, 8.97506925e-01, 8.50269719e-01],
[1.00000000e+00, 1.00000000e+00, 1.00000000e+00, 1.00000000e+00]]])
```


```
[18]: lin_reg_3 = LinearRegression().fit(X_poly_3, y)
```

```
[19]: X_line_poly_3 = poly.fit_transform(X_line.reshape(-1, 1))
plt.plot(X_line, lin_reg_3.predict(X_line_poly_3))
plot_sin()
```



```
[20]: # Check overfitting.
fig, axes = plt.subplots(1, 3, figsize=(16, 4))

for degree, ax in zip([5, 15, 25], axes):
    poly = PolynomialFeatures(degree=degree)
    X_poly = poly.fit_transform(X.reshape(-1, 1))
    lin_reg = LinearRegression().fit(X_poly, y)
    X_line_poly = poly.fit_transform(X_line.reshape(-1, 1))
    ax.plot(X_line, lin_reg.predict(X_line_poly))
    ax.scatter(X, y)
    ax.plot(X_line, sin_X, 'red')
```



1.3 Ridge and Lasso regressions

```
[21]: !pip install -q mglearn
import mglearn
import pandas as pd
from sklearn.model_selection import train_test_split
```

```
[22]: X, y = mglearn.datasets.load_extended_boston()
```

```
[23]: X
```

```
[23]: array([[0.00000000e+00, 1.80000000e-01, 6.78152493e-02, ...,
            1.00000000e+00, 8.96799117e-02, 8.04248656e-03],
            [2.35922539e-04, 0.00000000e+00, 2.42302053e-01, ...,
            1.00000000e+00, 2.04470199e-01, 4.18080621e-02],
            [2.35697744e-04, 0.00000000e+00, 2.42302053e-01, ...,
            9.79579831e-01, 6.28144504e-02, 4.02790570e-03],
            ...,
            [6.11892474e-04, 0.00000000e+00, 4.20454545e-01, ...,
            1.00000000e+00, 1.07891832e-01, 1.16406475e-02],
            [1.16072990e-03, 0.00000000e+00, 4.20454545e-01, ...,
            9.82676920e-01, 1.29930407e-01, 1.71795127e-02],
            [4.61841693e-04, 0.00000000e+00, 4.20454545e-01, ...,
            1.00000000e+00, 1.69701987e-01, 2.87987643e-02]])
```

```
[24]: y
```

```
[24]: array([24. , 21.6, 34.7, 33.4, 36.2, 28.7, 22.9, 27.1, 16.5, 18.9, 15. ,
            18.9, 21.7, 20.4, 18.2, 19.9, 23.1, 17.5, 20.2, 18.2, 13.6, 19.6,
            15.2, 14.5, 15.6, 13.9, 16.6, 14.8, 18.4, 21. , 12.7, 14.5, 13.2,
            13.1, 13.5, 18.9, 20. , 21. , 24.7, 30.8, 34.9, 26.6, 25.3, 24.7,
            21.2, 19.3, 20. , 16.6, 14.4, 19.4, 19.7, 20.5, 25. , 23.4, 18.9,
            35.4, 24.7, 31.6, 23.3, 19.6, 18.7, 16. , 22.2, 25. , 33. , 23.5,
            19.4, 22. , 17.4, 20.9, 24.2, 21.7, 22.8, 23.4, 24.1, 21.4, 20. ,
            20.8, 21.2, 20.3, 28. , 23.9, 24.8, 22.9, 23.9, 26.6, 22.5, 22.2,
```

```

23.6, 28.7, 22.6, 22. , 22.9, 25. , 20.6, 28.4, 21.4, 38.7, 43.8,
33.2, 27.5, 26.5, 18.6, 19.3, 20.1, 19.5, 19.5, 20.4, 19.8, 19.4,
21.7, 22.8, 18.8, 18.7, 18.5, 18.3, 21.2, 19.2, 20.4, 19.3, 22. ,
20.3, 20.5, 17.3, 18.8, 21.4, 15.7, 16.2, 18. , 14.3, 19.2, 19.6,
23. , 18.4, 15.6, 18.1, 17.4, 17.1, 13.3, 17.8, 14. , 14.4, 13.4,
15.6, 11.8, 13.8, 15.6, 14.6, 17.8, 15.4, 21.5, 19.6, 15.3, 19.4,
17. , 15.6, 13.1, 41.3, 24.3, 23.3, 27. , 50. , 50. , 50. , 22.7,
25. , 50. , 23.8, 23.8, 22.3, 17.4, 19.1, 23.1, 23.6, 22.6, 29.4,
23.2, 24.6, 29.9, 37.2, 39.8, 36.2, 37.9, 32.5, 26.4, 29.6, 50. ,
32. , 29.8, 34.9, 37. , 30.5, 36.4, 31.1, 29.1, 50. , 33.3, 30.3,
34.6, 34.9, 32.9, 24.1, 42.3, 48.5, 50. , 22.6, 24.4, 22.5, 24.4,
20. , 21.7, 19.3, 22.4, 28.1, 23.7, 25. , 23.3, 28.7, 21.5, 23. ,
26.7, 21.7, 27.5, 30.1, 44.8, 50. , 37.6, 31.6, 46.7, 31.5, 24.3,
31.7, 41.7, 48.3, 29. , 24. , 25.1, 31.5, 23.7, 23.3, 22. , 20.1,
22.2, 23.7, 17.6, 18.5, 24.3, 20.5, 24.5, 26.2, 24.4, 24.8, 29.6,
42.8, 21.9, 20.9, 44. , 50. , 36. , 30.1, 33.8, 43.1, 48.8, 31. ,
36.5, 22.8, 30.7, 50. , 43.5, 20.7, 21.1, 25.2, 24.4, 35.2, 32.4,
32. , 33.2, 33.1, 29.1, 35.1, 45.4, 35.4, 46. , 50. , 32.2, 22. ,
20.1, 23.2, 22.3, 24.8, 28.5, 37.3, 27.9, 23.9, 21.7, 28.6, 27.1,
20.3, 22.5, 29. , 24.8, 22. , 26.4, 33.1, 36.1, 28.4, 33.4, 28.2,
22.8, 20.3, 16.1, 22.1, 19.4, 21.6, 23.8, 16.2, 17.8, 19.8, 23.1,
21. , 23.8, 23.1, 20.4, 18.5, 25. , 24.6, 23. , 22.2, 19.3, 22.6,
19.8, 17.1, 19.4, 22.2, 20.7, 21.1, 19.5, 18.5, 20.6, 19. , 18.7,
32.7, 16.5, 23.9, 31.2, 17.5, 17.2, 23.1, 24.5, 26.6, 22.9, 24.1,
18.6, 30.1, 18.2, 20.6, 17.8, 21.7, 22.7, 22.6, 25. , 19.9, 20.8,
16.8, 21.9, 27.5, 21.9, 23.1, 50. , 50. , 50. , 50. , 50. , 13.8,
13.8, 15. , 13.9, 13.3, 13.1, 10.2, 10.4, 10.9, 11.3, 12.3, 8.8,
7.2, 10.5, 7.4, 10.2, 11.5, 15.1, 23.2, 9.7, 13.8, 12.7, 13.1,
12.5, 8.5, 5. , 6.3, 5.6, 7.2, 12.1, 8.3, 8.5, 5. , 11.9,
27.9, 17.2, 27.5, 15. , 17.2, 17.9, 16.3, 7. , 7.2, 7.5, 10.4,
8.8, 8.4, 16.7, 14.2, 20.8, 13.4, 11.7, 8.3, 10.2, 10.9, 11. ,
9.5, 14.5, 14.1, 16.1, 14.3, 11.7, 13.4, 9.6, 8.7, 8.4, 12.8,
10.5, 17.1, 18.4, 15.4, 10.8, 11.8, 14.9, 12.6, 14.1, 13. , 13.4,
15.2, 16.1, 17.8, 14.9, 14.1, 12.7, 13.5, 14.9, 20. , 16.4, 17.7,
19.5, 20.2, 21.4, 19.9, 19. , 19.1, 19.1, 20.1, 19.9, 19.6, 23.2,
29.8, 13.8, 13.3, 16.7, 12. , 14.6, 21.4, 23. , 23.7, 25. , 21.8,
20.6, 21.2, 19.1, 20.6, 15.2, 7. , 8.1, 13.6, 20.1, 21.8, 24.5,
23.1, 19.7, 18.3, 21.2, 17.5, 16.8, 22.4, 20.6, 23.9, 22. , 11.9])

```

```

[25]: df_X = pd.DataFrame(X)
      df_y = pd.DataFrame(y)

```

```

[26]: df_X

```

```

[26]:
      0      1      2      3      ...      100      101      102      103
0  0.000000  0.18  0.067815  0.0  ...  0.025759  1.000000  0.089680  0.008042
1  0.000236  0.00  0.242302  0.0  ...  0.113111  1.000000  0.204470  0.041808

```

2	0.000236	0.00	0.242302	0.0	...	0.035109	0.979580	0.062814	0.004028
3	0.000293	0.00	0.063050	0.0	...	0.021667	0.988585	0.033197	0.001115
4	0.000705	0.00	0.063050	0.0	...	0.064464	1.000000	0.099338	0.009868
..
501	0.000633	0.00	0.420455	0.0	...	0.195787	0.975392	0.216382	0.048003
502	0.000438	0.00	0.420455	0.0	...	0.181239	1.000000	0.202815	0.041134
503	0.000612	0.00	0.420455	0.0	...	0.096414	1.000000	0.107892	0.011641
504	0.001161	0.00	0.420455	0.0	...	0.117127	0.982677	0.129930	0.017180
505	0.000462	0.00	0.420455	0.0	...	0.151649	1.000000	0.169702	0.028799

[506 rows x 104 columns]

```
[27]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

```
[28]: lin_reg = LinearRegression().fit(X_train, y_train)
```

```
[29]: def print_score(model):
        print(round(model.score(X_train, y_train), 3))
        print(round(model.score(X_test, y_test), 3))
```

```
[30]: print_score(lin_reg)
```

0.952

0.607

```
[31]: from sklearn.linear_model import Ridge, Lasso
```

1.4 Ridge regressions

```
[32]: ridge = Ridge().fit(X_train, y_train)
```

```
[33]: print_score(ridge)
```

0.886

0.753

```
[34]: ridge_10 = Ridge(alpha=10).fit(X_train, y_train)
        print_score(ridge_10)
```

0.788

0.636

```
[35]: ridge_01 = Ridge(alpha=0.1).fit(X_train, y_train)
        print_score(ridge_01)
```

0.928

0.772


```
[36]: coefficients = pd.DataFrame({'lin_reg': lin_reg.coef_, 'ridge':ridge.coef_,
    ↪ 'ridge_10':ridge_10.coef_, 'ridge_01':ridge_01.coef_})
```

```
[37]: coefficients
```

```
[37]:
```

	lin_reg	ridge	ridge_10	ridge_01
0	-412.710947	-1.413684	-0.811610	-2.333007
1	-52.243207	-1.556619	0.647609	-5.441929
2	-131.898815	-1.465434	-0.809783	-6.128655
3	-12.004136	-0.126616	0.311335	-0.778525
4	-15.510713	-0.079119	-0.684624	1.291530
..
99	35.361714	-2.361124	-1.769550	-0.559662
100	11.955343	0.043467	-2.036371	4.947479
101	0.677026	1.208860	1.079849	2.667765
102	2.734520	-6.325992	-3.952641	-6.907626
103	30.372001	10.360023	0.142449	24.258596

```
[104 rows x 4 columns]
```

1.5 Lasso regressions

```
[38]: lasso = Lasso().fit(X_train, y_train)
```

```
[39]: print_score(lasso)
```

```
0.293
```

```
0.209
```

```
[40]: lasso_001 = Lasso(alpha=0.01, max_iter=10000).fit(X_train, y_train)
print_score(lasso_001)
```

```
0.896
```

```
0.766
```

```
[41]: coefficients_lasso = pd.DataFrame({'lin_reg':lin_reg.coef_, 'lasso':lasso.
    ↪coef_, 'lasso_001':lasso_001.coef_})
```

```
[42]: coefficients_lasso
```

```
[42]:
```

	lin_reg	lasso	lasso_001
0	-412.710947	-0.0	-0.000000
1	-52.243207	0.0	-0.000000
2	-131.898815	-0.0	-0.000000
3	-12.004136	0.0	0.000000
4	-15.510713	-0.0	-0.000000
..

99	35.361714	-0.0	0.000000
100	11.955343	-0.0	-0.000000
101	0.677026	0.0	0.344041
102	2.734520	-0.0	-8.246456
103	30.372001	-0.0	17.560672

[104 rows x 3 columns]